

The Disclosure Verifier

An agentic system that checks whether a company's own words match its own numbers — extracting quantitative claims from SEC filing prose and reconciling each one against the structured data the company itself reported, **with a citation back to the exact filing.**

Status Phases 0–7 complete **Tests** 260 passing, CI green

Repo github.com/asim-aa/disclosure-verifier

"Revenue for fiscal year 2026 was \$215.9 billion, up 65% from a year ago."

extracted → {metric: Revenue, value: 215,900,000,000, type: absolute}

extracted → {metric: Revenue, value: 65.0, type: growth_pct}

reconciled → **CONSISTENT** EDGAR reports \$215,938,000,000 – 0.02% off

reconciled → **INCONSISTENT** EDGAR implies 46.09% growth, not 65%

01

DSPy optimization, measured against a baseline

The capstone's Pillar 2 requirement: at least one prompt signature optimized with DSPy, scored against a hand-written baseline — not just used, *proven*. Ground truth is 78 examples / 161 claims / 9 true negatives, hand-labeled from real MSFT and NVDA filings before any DSPy code was written, to keep the measurement honest.

EXTRACTOR	PRECISION	RECALL	F1
Hand-written baseline prompt	0.711	0.750	0.730
DSPy zero-shot (ChainOfThought, no demos)	0.763	0.806	0.784
DSPy optimized (BootstrapFewShot)	0.763	0.806	0.784

Held-out test set, 16 examples never seen during optimization. Delta vs. baseline:
+0.054 F1

The honest framing matters more than the win — including a result that didn't survive its own honesty check. Switching the signature from `Predict` to `ChainOfThought` — reasoning before committing to structured output — took zero-shot DSPy from underperforming the baseline (an earlier run measured 0.676 F1 on the bare `Predict` signature) to clearly beating it (0.784), before any optimization ran. `BootstrapFewShot` optimization on top of that produced *exactly identical* pooled numbers: same precision, same recall, the same 29/9/7 true-positive/false-positive/false-negative split. Verified this wasn't an evaluation bug — the compiled program is a genuinely separate instance, and the optimizer log confirms 4 real demonstrations were bootstrapped.

CHECKED AGAINST ITS OWN NOISE FLOOR

A 16-example held-out set produces 45 claim-level tp/fp/fn decisions — small enough that `eval/run_comparison.py` now computes the sampling noise floor ($SE(p) = \sqrt{p(1-p)/n}$, ~95% half-width $1.96 \cdot SE$) alongside every delta, rather than letting a modest F1 gain read as a settled result. At n=45, that half-width is ± 0.120 — wider than the $+0.054$ delta above. The pooled win is not distinguishable from sampling noise at this test set size; a bigger held-out set is what it would take to resolve this, not a cleaner summary sentence.

A stratified, per- `comparison_type` breakdown adds a second reason the pooled number understates the picture: optimization moved the two claim categories in

opposite directions. `absolute` claims improved ($0.625 \rightarrow 0.788$ F1), while `absolute_change` claims got worse ($0.429 \rightarrow 0.308$) — a regression a pooled average hides because `growth_pct` claims (1.000 F1 throughout) dominate the sample and keep the aggregate looking flat-to-positive. Reasoning mattered more than optimization for this task, and the optimization step itself is a wash at best and a net negative for one category at worst — a more interesting and more honest finding than a clean "optimization wins" story would have been.

TRIED THE NEXT OPTIMIZER UP, HONESTLY

`BootstrapFewShot` adding nothing was the reason to try `dspy.GEPA` — a reflective optimizer that rewrites the instruction itself, guided by *why* specific attempts failed rather than a bare pass/fail score. Caveat stated before the result, not after: GEPA's reflection step is meant to be guided by a model stronger than the one being optimized, and this project has exactly one LLM endpoint (gpt-oss-20b) — it reflects on its own failures here, not a stronger model's. Real run, `auto="light"`, 444 metric calls, 37 minutes on the local single-GPU server: F1 $0.784 \rightarrow 0.800$, a $+0.016$ delta *inside* the ± 0.117 noise floor at this sample size — not distinguishable from chance, same honest conclusion as `BootstrapFewShot`. The per-category read is more telling than the pooled number: `growth_pct` reached a perfect 1.000 F1 and `absolute` improved to 0.857, but `absolute_change` collapsed to 0.167 — a trade-off, not a clean win.

THE REPEATED-TOUCH CAVEAT, FINALLY CLOSED

Every number above has one honest asterisk: the same 16-example test slice was scored on every run — baseline, zero-shot, `BootstrapFewShot`, `GEPA` — so every delta was "the best estimate from a set we've looked at many times," never a clean holdout result. Closed directly: 25 fresh examples / 59 claims, hand-labeled by reading real AMZN and AAPL 10-K MD&A text — two companies absent from the training data entirely, scored for the first time by this run. The result actually *moves*: baseline collapsed on genuinely new companies ($0.730 \rightarrow 0.344$ F1, 44 false positives in one category), while zero-shot held up ($0.784 \rightarrow 0.752$). DSPy vs. baseline is now a real, **noise-floor-clearing** result — $+0.355$ F1 at $n=82$, noise floor ± 0.099 — resolved for the first time instead of perpetually "inconclusive at this sample size." A new question opened with it: `BootstrapFewShot` didn't just fail to help on the fresh set, it measurably *hurt* ($0.752 \rightarrow 0.698$) — consistent with its 4 demos overfitting to MSFT/NVDA's phrasing, though at $n=25$ that specific drop is itself still inside its own noise floor. One finding closed, a sharper one opened — not a tidier story, a truer one.

02

RLVR/GRPO fine-tuning, checked the same way

The capstone's Pillar 4 requirement: reinforcement learning from verifiable rewards — training a small open model to reason through the arithmetic the Numerical Reconciler already does deterministically, using the Reconciler itself as the reward, audited first (15/15 correct, zero false-"consistent" results — see finding 08 below).

Qwen2.5-7B-Instruct , QLoRA rank 32 via Unsloth, a single RTX 4090. Same discipline as Pillar 2: baseline vs. trained, checked against the noise floor, not just reported as a bare delta.

	ACCURACY	MEAN REWARD	FALSE-CONSISTENT
Baseline (zero-shot)	0.766	0.633	0.148
Trained (1 epoch, 1,300 GRPO steps)	0.855	0.826	0.036

Held-out test set, 304 examples. Delta vs. baseline: **+0.089 accuracy**

Unlike Pillar 2's optimization delta, this one survives its own noise-floor check: at n=304 the combined 95% half-width is ± 0.062 — narrower than the +0.089 delta. The false-consistent rate — the dangerous failure mode, a model confidently endorsing a wrong claim — dropped from 0.148 to 0.036, also clearing its own noise floor (± 0.045 against a -0.112 delta). Format failures (no parseable verdict) went from 1/304 to 0/304.

COMPARISON_TYPE	BASELINE	TRAINED
absolute	0.949	0.975
absolute_change	0.589	0.821
bps_change	0.233	0.400
growth_pct	0.810	0.870

absolute_change shows the largest per-category gain, also noise-floor-clearing at this stratum's n=56. bps_change — two ratios, then a basis-point difference, the hardest task in the set — moved from 7/30 to 12/30 correct, but at n=30 that's not provable on its own; a first, partial run (150 steps, ~31 minutes) had left it at an exact, unchanged null. Directionally real, not yet statistically defensible — stated honestly rather than rounded up.

Same held-out example, both models, greedy decoding — reasoning *quality*, not a verdict flip (both reach the correct answer here):

BASE MODEL

Step 1: Calculate the actual change in revenue.

Actual change = Current period value – Prior period value
 = 274,515,000,000 USD – 2260,174,000,000 USD
 = 14,341,000,000 USD

Step 2: Check if the claimed value falls within the tolerance range.

Tolerance = $0.01 * \max(2260,174,000,000 \text{ USD}, 274,515,000,000 \text{ USD})$
 = 22,701,500,000 USD

Step 3: Compare the claimed value to the tolerance range.

Tolerance range: –22,701,500,000 USD to 22,701,500,000 USD

Since the claimed value falls within the tolerance range, the claim is consistent.

VERDICT: consistent

TRAINED MODEL (1,300 STEPS)

Step 1: Identify the relevant values from the reported data.

- Revenue for period ending 2020-09-26: 274,515,000,000 USD
- Revenue for period ending 2019-09-28: 260,174,000,000 USD

Step 2: Calculate the actual change in revenue.

Actual change = 274,515,000,000 USD – 260,174,000,000 USD
 = 14,341,000,000 USD

Step 3: Compare the claimed value with the calculated actual change.

Claimed value: 14,341,000,000 USD | Actual change: 14,341,000,000 USD

The claimed value matches the actual calculated change.

VERDICT: consistent

The base model's arithmetic typos a number mid-calculation (2260,174,000,000 instead of 260,174,000,000) yet still lands on the right difference — the final number reads recalled, not derived. It also checks the claim against a symmetric tolerance band around zero, a different and weaker test than what it's actually supposed to check, and only passes here because the claimed change happens to be small relative to that band. The trained model computes the actual change cleanly and compares it directly to the claim — the same check the real Reconciler performs. That's the shape of improvement behind the numbers above.

WHAT IT TOOK TO GET HERE

Three real upstream packaging bugs, none about the reward design itself: TRL's `GRPOTrainer` unconditionally imports an unmaintained, incompatible `llm_blender` and hits a `weave` availability check that false-positives in this environment, for judge/logging features never used here — worked around with a `sys.modules` stub rather than chasing dependency versions for functionality the project doesn't touch. It also assumes a `model.warnings_issued` attribute the PEFT/Unsloth-wrapped model doesn't initialize — fixed with a one-line defensive guard. Full writeup, including the first (inconclusive, honestly reported) 150-step run: `docs/phase7-results.md`.

RETRAINED ON THE FIXED DATASET – THE HONEST RESULT

Finding 09 below traces `bps_change`'s weak performance to a data-generator bug, not a reasoning limitation — the dataset had only 2 real distinct ratios to draw from. Fixed and retrained for real on the same box: 3,403 examples (up from 1,612), same architecture, same 1,300 GRPO steps, evaluated on a fresh 706-example held-out split (not comparable row-for-row to the 304-example table above — a different split of a larger dataset).

	BASE	TRAINED	DELTA	NOISE FLOOR
Overall accuracy	0.820	0.858	+0.038	±0.038
False-consistent rate	0.108	0.020	−0.088	±0.025
absolute	0.969	0.994	+0.025	±0.021
absolute_change	0.667	0.825	+0.158	±0.108
growth_pct	0.858	0.893	+0.035	±0.070
bps_change	0.436	0.372	−0.064	±0.140

Held-out test set, n=706. Noise floor: combined 95% half-width across the two proportions.

The false-consistent rate improved again, by a wide, noise-floor-clearing margin (0.108 → 0.020) — the dangerous failure mode kept shrinking on a second, independent run. `absolute_change` improved substantially and really (+0.158). But the honest headline is `bps_change` itself: the fix that tripled its raw training-example count (114 → 350) produced **no measurable win** — the number actually moved

down slightly ($0.436 \rightarrow 0.372$), though well inside the noise floor at this small stratum ($n=94$), so not proven worse either, just not proven better. A plausible reason, not yet confirmed: `max_steps` stayed fixed at 1,300 while the dataset grew $2.1\times$ — so despite the raw `bps_change` count tripling, every example (that category included) got proportionally *less* repeated exposure than in the original run. If that's right, the next real experiment is more steps on the bigger dataset, not the reportioned data alone — flagged as follow-up work, not run here.

TESTED THE STEP-BUDGET HYPOTHESIS DIRECTLY — IT DIDN'T HOLD

Ran the follow-up flagged above: same dataset, same architecture, `max_steps` scaled from 1,300 to **2,730** to match the dataset's $2.1\times$ growth, same 706-example held-out split. Result: overall accuracy $0.820 \rightarrow 0.865$ and false-consistent rate $0.108 \rightarrow 0.024$, both real again; `absolute_change` reached $0.667 \rightarrow 0.867$ — the largest clean win across all three runs. But `bps_change` moved to only 0.404 — numerically closer to the 0.436 baseline than the 1,300-step run's 0.372, but that move is itself inside the noise floor against the 1,300-step result (± 0.139), so not a real change either. Doubling the steps didn't fix it. The step-budget-dilution hypothesis is the one that turned out to be wrong here — whatever's actually holding `bps_change` back isn't simply a matter of more repeated exposure, reported as the real, if less satisfying, answer rather than left as an open theory that sounded plausible.

03

A real run, at scale

Not a cached demo — this is the actual coordinator, run against NVDA's FY2026 10-K, live EDGAR data and a live LLM call for every extraction.

NVDA · FY2026 10-K · MD&A		13 claims verified
CONSISTENT	Revenue EDGAR: \$215,938,000,000 — 0.02% off	\$215.9B
INCONSISTENT	Revenue growth EDGAR implies 46.09% growth, not the stated figure	65.0%
CONSISTENT	Income tax expense EDGAR: \$21,383,000,000 — 0.08% off	\$21.4B
UNVERIFIABLE	Data Center revenue growth segment-level — no top-level XBRL tag exists; declined rather than guessed	68.0%

13 claims checked	22 tool calls	1.39s wall-clock	0 crashes
----------------------	------------------	---------------------	--------------

One filing is a worked example, not a scale claim. Phase 6 runs the same real coordinator — no mocks — against 5 companies' real 10-Ks, including GOOGL and AMZN, deliberately never used while building any of this, to check whether the pipeline generalizes or was quietly tuned to the three it had already seen.

5/5 tickers succeeded	18 claims verified	128 tool calls	0 crashes
--------------------------	-----------------------	-------------------	--------------

Budget-capped per ticker (25 chunks / 20 extraction calls / 240s) — a real 10-K can run 100+ MD&A chunks, each an LLM call.

The first run surfaced two real coverage gaps rather than a clean pass — the kind of thing three hand-picked companies can hide. GOOGL and AMZN returned *zero* claims: their real body heading for "Item 7. Management's Discussion..." renders as an isolated "Item 7." node with the title following separately, structurally identical

to what the parser had been treating as a table-of-contents-only pattern — a heuristic that looked solid against 3 companies and broke on the 4th. Fixed by scoring every candidate heading window by the gap to its nearest end-marker instead of by which text node it landed in (a real section runs hundreds of lines; a ToC entry is a few). Separately, most real MD&A prose past the top-line numbers is segment-specific ("Azure revenue," "H2O charge," "XBOX content and services revenue") — investigated rather than assumed fixable: SEC's company-facts API returns every data point as a flat record with no segment/member/axis field at all, so those claims correctly decline rather than guess. But not everything in that unresolved list was actually segment-level — "Commercial remaining performance obligation" turned out to map to `RevenueRemainingPerformanceObligation`, a real top-level concept just missing from the dictionary. Re-run after both fixes: MSFT's remaining-performance-obligation claims now resolve to real verdicts instead of "unverifiable" — **CONSISTENT** on the absolute figure (0.88% off), **INCONSISTENT** on the growth figure, itself flagged as possibly a known period-matching limitation rather than a real miss.

THE OTHER HARNESS CLAIM, EXERCISED LIVE TOO

Budget enforcement shows up in every run above (every ticker hit its cap).

Checkpoint/resume — the other half of the harness's real-conditions claim — had only ever been proven against mock scenarios. A dedicated run against MSFT's real 215-chunk filing closes that: a tiny budget forces a stop at chunk 3, a real checkpoint lands on disk, a resume loads it, skips re-extracting those 3 chunks, and continues forward to chunk 23 with 11 real claims accumulated — the save → load → skip → continue cycle, proven against a live filing, not a fixture.

04 What actually broke

The interesting engineering here wasn't the happy path — it was what real SEC data did to the happy path. Each of these was caught by testing against live data, not by trusting the design on paper.

01 A metric-matching shortcut that would verify the wrong number

Resolving a claim's free-text metric ("Azure revenue") to an exact XBRL concept originally tolerated wording variance via substring matching. But "revenue" is a substring of "Azure and other cloud services revenue" — so a segment's revenue would have been silently checked against *total company* revenue and returned a confidently wrong "consistent" verdict. Fixed to exact-match plus explicit aliases only.

[agents/resolver.py](#)

02 XBRL doesn't tag percentages as percentages

Rate concepts like effective tax rate are reported with `unit="pure"` — a decimal fraction (`0.19` , not `19`). A claim stating "19%" would have failed to match the fact at all, and once that was fixed, would have compared `19.0` against `0.19` and read as wildly wrong even when correct.

[agents/verification_agent.py](#)

03 Companies retag their own concepts mid-history

NVDA reported revenue under

`RevenueFromContractWithCustomerExcludingAssessedTax` through FY2022, then switched to plain `Revenues` . The old tag never disappears from the company's concept list — it just stops receiving new data. Picking "the first candidate that exists at all" silently locked onto a tag with no data newer than 2022. Fixed to pick by data recency instead.

[agents/resolver.py](#)

04

A newer filing can corrupt an older one's claims

A claim extracted from a 10-K's MD&A describes *that 10-K's* fiscal year — but a newer 10-Q, almost always already filed by the time verification runs, has a more recent quarterly `period_end`. Without a cutoff, that quarterly figure would get picked as "current," comparing the wrong two numbers entirely. Fixed with an `as_of` anchor tied to the claim's own source filing.

[agents/resolver.py](#)

05

A bug in the measurement itself, caught before it could lie

An early precision/recall scorer would have counted a generic prediction of `"Revenue"` as matching gold label `"Microsoft Cloud revenue"` — a bug in the *grader*, which is worse than a bug in the thing being graded, because it makes bad results look good. Its own unit test caught it before the real comparison ever ran.

[eval/metrics.py](#)

06

Reasoning silently corrupted its own numeric output

Switching the extraction signature from `dspy.Predict` to `dspy.ChainOfThought` made the model write dollar amounts in shorthand inside its reasoning — "\$215.9 billion" — then just echo that literal `215.9` into the structured `value` field instead of expanding it to `215900000000`. Consistent, every run. Plain `Predict` never had this failure mode, because it never generates that intervening shorthand text to anchor on. Fixed by instructing the signature to write the fully-expanded number in the reasoning itself, so the structured step has nothing left to get wrong.

[eval/dspy_extractor.py](#)

07

A restatement cutoff that resolved periods but not values

The verification agent already anchored *period and concept* resolution to an `as_of` cutoff — the claim's own source filing date — so a 10-Q couldn't leak into a 10-K's period selection (finding 04). But that same cutoff never reached the reconciler's own fact lookup: a claim from an older filing could still be compared against a company's *later restatement* of that same period and get flagged inconsistent, even though it was accurate when made. Fixed by threading `as_of` through `reconcile()` itself, with a regression test built from a live-numbers reproduction of the bug.

[tools/reconciler.py](#)

08

Auditing the reward function before it has a job

The Reconciler is the ground-truth signal for Phase 7's RLVR/GRPO reward — which means an exploitable seam in it wouldn't just add noise, it would give policy optimization something specific to find and amplify. Before any GPU time is spent, `eval/reconciler_audit.py` now runs 15 known-good, known-bad, and adversarial cases (sign flips, magnitude confusion, comparison-type mislabeling, tolerance-boundary probes) as a permanent CI check. Result: 15/15 correct verdicts, zero false-"consistent" outcomes — the one failure mode worth watching for, since a false alarm is costly but a missed inconsistency is the dangerous direction.

[eval/reconciler_audit.py](#)

09

A weak category traced to the data generator, not the model

`bps_change` stayed the worst RLVR category through training (0.233 → 0.400) — the easy read is "the hardest reasoning task, needs more steps." The actual cause was upstream: the dataset generator only had *2 real distinct ratios* to draw from (gross margin, operating margin — 2 of its original 4 "pairs" were the same ratio under a company's alternate revenue tag, not a second ratio), which made `bps_change` 8.7% of training data against 43.3% for `absolute`. Not rare reasoning — rare *examples*. Fixed by adding 5 more ratio pairs (net margin, R&D intensity, SG&A ratio, cost ratio, opex ratio) from concepts already verified present in real company data: `bps_change`'s share rose to 13.0%, raw count 114 → 350. Retrained on it for real (\$02) — the honest result is a null, not a win: `bps_change` didn't clear its own noise floor either direction, most plausibly because the GRPO step budget didn't scale with the larger dataset. A cleanly traced root cause, fixed, and still not sufficient on its own — the kind of finding worth keeping, not smoothing over.

[phase7/build_dataset.py](#)

10

A claim's own stated period was being ignored entirely

`resolve_periods()` always picked the globally most-recent fact as "current," regardless of what the claim's own period text actually said — so two claims from the same sentence naming two different fiscal years ("Income tax expense was \$21.4 billion and \$11.1 billion for fiscal years 2026 and 2025, respectively") both resolved to the *same* (current, comparison) pair, and the FY2025 claim ended up checked against the FY2026 fact. Confirmed against this exact live NVDA sentence: without a hint, the FY2025 claim's "current" resolved to the FY2026 annual figure (\$21.38B) against a \$13.9B quarterly comparison — nonsensical. Fixed by threading the claim's own period text through as a hint: an explicit fiscal year now picks that year's fact as current, and "sequentially" vs. "a year ago" now correctly distinguish the immediately-preceding period from the same period ~12 months back. Re-run against the same live sentence: both claims now independently resolve **CONSISTENT** against their own correct periods.

[agents/resolver.py](#)

11

A quarter reference resolved to the wrong duration entirely

A named ordinal quarter ("the third quarter", "Q3") had been deliberately left unresolved — real filings also name quarters by calendar month ("the September quarter"), and guessing that mapping wrong without knowing a company's fiscal-year-end risks a wrong-but-confident match. The ordinal form is safe, though: XBRL's own `fiscal_period` field is itself fiscal-quarter-numbered, so "the third quarter" maps directly to `fp="Q3"` with no guessing involved. Implementing it surfaced a second, sharper bug before it shipped: a 10-Q commonly tags *both* the standalone 3-month figure and a 6-/9-month year-to-date cumulative with the identical `fiscal_period` and `period_end` — an unfiltered match against real NVDA data returned a \$91.166B "Q3" figure that was actually the 9-month cumulative, not the \$35.082B standalone quarter. Fixed by preferring the ~91-day-length fact when both exist. Calendar-named quarters are still correctly declined rather than guessed.

[agents/resolver.py](#)

OPEN, NOT HIDDEN

A quarter named by calendar month ("the September quarter", "the December quarter") still isn't parsed into a specific fiscal period — translating that safely needs the company's own fiscal-year-end, which isn't available at this resolution step, so it still falls back to the default "most recent" pick rather than guessing. Flagged in [agents/resolver.py](#) as follow-up, not silently left broken.

05

A real backtest: would this have caught an actual restatement?

Every result above is checked against real filings, but always in the present tense — does the pipeline verify a claim correctly *today*? A sharper question: has this project's own bitemporal machinery ever actually caught something real happening to a real company? SEC 8-K "Item 4.02" disclosures — a company telling investors its own past financial statements can no longer be relied on — supply real ground truth to test that against, using nothing synthetic.

13,827 real restatement fingerprints found	541 companies, from real 8-K filings	122 matched to real MD&A prose, 16 companies	118/122 flagged inconsistent — 96.7%
--	--	--	--

Three stages, each verified before moving to the next. **Finding restatements:** SEC's own XBRL company-facts data keeps both the original and the corrected value for the same (concept, period) when a company amends a filing — group by (concept, period), keep only groups where a 10-K/A or 10-Q/A specifically reports a different value. A looser first pass (any later-filed value counts) found 42,200 "fingerprints" that fell apart on inspection — 82% were routine reclassifications, not corrections. Requiring the later value come from an actual amendment dropped that to a defensible 13,827, across 541 companies. **Matching to prose:** for the top 60 of those companies by fingerprint magnitude, fetch the *original* filing's MD&A and run the project's real extraction agent over it, keeping claims whose metric resolves to the same concept and whose value is within 5% of the fingerprint — prose commonly rounds ("\$.1.8 billion" for \$1,838,000,000). 409 filings tried, 122 real claims matched this way, across 16 companies — Discover Financial Services and Rithm Capital Corp. (the pair the reliability fix below was first proven against) plus AZEK, Seneca Foods, Camber Energy, Astrana Health, Sanmina, Compass Minerals, and nine more. **The backtest itself:** for each matched claim, run the actual `reconcile()` function — the same code every other result in this document uses — twice: once with the bitemporal `as_of` cutoff set to the claim's own filing date (what did the reconciler know when the claim was made?), once set to the restated filing's date (what does it know now?). No LLM calls in this step — pure arithmetic against already-known claim values.

A sample of the 122 matched claims, illustrating all outcomes			as filed → after restatement
→ FAIL	Discover Financial Services — net income, Q3 2023 consistent at filing (0.0% off) → inconsistent once restated (16.55% off)		\$683M claimed
→ FAIL	Camber Energy — net income, Q2 2023 consistent at filing (0.0% off) → inconsistent once restated (250.53% off — one of the more dramatic corrections in the set)		\$1.92M claimed
FAIL → FAIL	Discover Financial Services — net income, Q1 2022 already inconsistent at filing (3.38% off) — the prose's own rounding missed tolerance before any restatement existed		\$1.2B claimed
PASS → PASS	Discover Financial Services — net income, FY2021 stayed consistent even after restatement — the 1.12% correction was smaller than the prose's own rounding, both sides land inside the 1% tolerance		\$5.4B claimed

102 of 122 show the clean target pattern — consistent when the claim was made, inconsistent once the restatement existed, using only data that existed at each point in time. 16 more were already inconsistent at the original filing, for an unrelated reason: the prose's own rounding fell outside tolerance regardless of any later restatement — still a real catch, just not evidence the bitemporal mechanism specifically did the work. 3 are structurally unverifiable on both sides — `absolute_change` claims ("decreased \$33.4 million") where Phase B's matching step doesn't capture the comparison period this claim type needs, so `reconcile()` correctly declines rather than guessing one. 1 of 122 is a genuine miss, reported rather than hidden: a real 1.12% restatement that happened to sit inside the gap between the prose's rounding and both the original and corrected figures. Combined, 118 of the 119 checkable matches (99.2%) — 118 of all 122 (96.7%) — would have been flagged inconsistent by the time each restatement existed — not a synthetic eval number, a real outcome across 16 real companies' real regulatory history.

SCOPED, NOT EXHAUSTIVE

122 matches come from 16 companies out of the 541 with real restatement fingerprints — the top 60 by fingerprint magnitude, not the full population. A first full-scan attempt across those 60 stalled past two hours with no incremental saving, a real reliability bug fixed with hard per-chunk timeouts and a save after every filing; the fix was proven on 2 companies first, then run for real across all 60 (moved to an always-on box partway through, since the run outlives a laptop staying open). 13,827 fingerprints and 541 companies are counted directly from real SEC data; 122 is the number actually pushed all the way through prose-matching and backtesting, stated as what it is rather than rounded up to the larger pool it was drawn from.

06

Why build a deterministic reconciler instead of just asking an LLM?

Every result above proves this *pipeline* works. This tests whether the *architecture choice* did — structured extraction, exact concept resolution, then deterministic bitemporal arithmetic, instead of the obvious alternative: hand a general LLM the claim and the raw numbers and let it decide. Tested directly against two already-verified test sets, same LLM endpoint used everywhere else in this project, no new hand-labeling.

	DETERMINISTIC RECONCILER	NAIVE LLM BASELINE
Adversarial audit (15 cases)	15/15 (100%)	11/15 (73.3%)
Backtest, correct flip pattern	102/122 (83.6%)	50/122 (41.0%)
False-positive rate at original filing	0% (by construction)	49% (49/100)

On `eval/reconciler_audit.py` 's 15 adversarial cases (each already 100% correct for the deterministic reconciler, by construction), the naive baseline is handed the exact same facts and asked to judge via reasoning alone: **11/15 (73.3%)**. The miss worth sitting with: a claim off by **1,000,000×** (stated in millions against facts reported in raw dollars) read as **CONSISTENT** to the naive baseline — the exact dangerous failure mode this whole project is built around, and one plain division in `tools/reconciler.py` never gets wrong.

The sharper test: the 122-match real backtest, checking whether a naive read can do the bitemporal reasoning the `as_of` design exists for. Given the *raw, unfiltered* fact history — every filed value, not the one bitemporally-correct value `_find_fact()` would pick — and told the claim's filing date in plain English, the naive baseline reproduced the correct "consistent when made, inconsistent once restated" pattern in only **50 of the 102 cases (49%)** the deterministic pipeline got right. Worse, and more precise: of the 100 real cases that were genuinely consistent as of their own original filing, the naive baseline **falsely called 49 of them "inconsistent"** — bleeding the restatement's later information backward into a judgment about the past. Telling a model "only use data filed on or before this date" in a prompt does not reliably make it do that; writing the cutoff into the fact-lookup code does — the same lesson this project's own development history already learned once, in finding 07.

HONEST SCOPE

Run against this project's own self-hosted `gpt-oss-20b` endpoint, held constant to isolate "does structure help" from "is this specific model good at this" — a stronger commercial model would plausibly narrow the gap, especially on raw arithmetic, and is a real follow-up not run here. One prompt attempt per case, no retries or self-consistency voting. A real harness bug was caught and fixed while building this too: the first pass omitted the denominator concept's facts for `bps_change` cases, so the naive baseline correctly said "unverifiable" on data it was never shown — fixed before the numbers above were reported.

07

Does it cry wolf on clean filings?

The backtest only tests sensitivity — it shows the reconciler flags claims that turned out to be wrong, on companies that actually restated. Never checked the other way: run against companies that *didn't* restate, does it stay quiet on genuinely accurate claims? The real Coordinator, no mocks, against 8 large tech companies confirmed to carry no restatement fingerprint anywhere in Phase A's scan.

152 real claims, 8 companies	33 consistent	0 inconsistent	0/33 reconciler false positives, given correct inputs
--	-------------------------------------	--------------------------------------	---

An apparent 25% false-positive rate among the 24 resolved verdicts — worth investigating before reporting, not reporting at face value. It split into two real, distinct causes. The first pass actually found **14** inconsistent claims (58.3%). Digging into the largest cluster found a genuine, previously undiscovered bug: `resolve_periods()` 's *default* comparison-period pick (no explicit hint) had no check that the comparison period's duration matched the current one's. Confirmed against real TXN data — the company's own MD&A states plainly *"Revenue of \$17.68 billion increased \$2.04 billion, or 13.0%, ... compared to fiscal 2024"*, an accurate claim — but the resolver picked a Q3 2025 10-Q's **9-month year-to-date** fact (\$13.259B) over the correct FY2024 annual figure (\$15.641B), because the YTD fact's `period_end` sorted more recently. **Fixed**, with regression tests from the real numbers: 8 of the 14 original false positives resolved correctly afterward.

The remaining 6 looked, at first, like an extraction bug. Three TXN claims showed the identical pattern — a sentence shaped "[current value] ... compared with [prior value]" where the number checked against EDGAR was the second one, not the first:

QUOTE	CLAIMED	REAL (CURRENT PERIOD)
"Net income was \$5.00 billion compared with \$4.80 billion."	\$4.80B	\$5.001B
"...provision for income taxes was \$709 million compared with \$654 million."	\$654M	\$709M
"...effective tax rate...was 12.4% in 2025 compared with 12.0% in 2024."	12.0%	12.4%

The obvious read — extraction grabbing the wrong number — was wrong, and worth showing rather than quietly fixing in place. Calling `eval/dspy_extractor.py` directly against these exact TXN paragraphs found it already produced the correct claims every time: two separate `absolute` claims, right values, right order. Extraction was never the problem.

The real bug was in `agents/verification_agent.py`: each extracted claim is verified *independently*. Neither value in "\$5.00 billion compared with \$4.80 billion" states an explicit period, so both claims call `resolve_periods` with an empty hint and both land on the *same* most-recent fact — the \$5.00B claim matches it, the accurate \$4.80B prior-year claim doesn't, and gets flagged inconsistent for describing the wrong year on purpose. **Fixed** in `agents/coordinator.py`: repeated `(metric, value_unit)` claims within one chunk are numbered, and the second same-metric absolute claim with no stated period is checked against the comparison fact the first occurrence already resolved, instead of being re-resolved against current again. Keying by value unit rather than the claim's quoted text mattered in practice — a real TXN sentence ("Operating profit was \$6.02 billion, or 34.1% of revenue, compared with \$5.47 billion, or 34.9% of revenue.") produces four claims for one metric, and the model doesn't always give paired claims an identical quote span.

Re-verifying surfaced two more real, smaller bugs: a percent claim ("gross profit margin") mapped to a concept XBRL reports as a raw dollar figure, never a ratio — guaranteed to "fail" no matter the value, now correctly returned as unverifiable instead; and a bare year ("2024", no "fiscal"/"FY" prefix) that the period-hint parser didn't recognize, silently falling back to the wrong default period. Both fixed the same way as everything else here — confirmed against the live data first, then a tightly-scoped regression test.

Narrowing the remaining 9 misses (7 CRM, 2 CSCO) found two more real bugs — both period-resolution gaps, not extraction or chunking. First: a relative hint like "a year ago" describing an *absolute* claim's entire value, not a growth claim's delta reference. CRM's real MD&A states "*as compared to diluted net income per share of \$6.36 from a year ago*" as its own standalone chunk — the current-period sentence lives elsewhere, so there's no sibling claim to pair it with. `resolve_periods` already resolves the correct year-ago fact as its comparison period (built for exactly this phrase on growth claims) — the fix just checks an absolute claim whose own period matches that phrasing against the comparison fact instead of current. Confirmed it doesn't misfire on the legitimate use of the same words: CRM's own remaining-

performance-obligation claim ("increase of 14 percent *year-over-year*") is a real growth claim where the current period genuinely is now.

Second: a quarter hint with nothing to match. CSCO's real MD&A states *"For the fourth quarter of fiscal 2025... total revenue increased by 8%"* — but CSCO doesn't separately tag a standalone Q4 fact (it folds Q4 into the annual 10-K), so the hint matched nothing and silently fell back to comparing against full-year totals, flagging a real ~8% quarterly claim against ~5.3% annual growth. **Fixed** by raising instead of guessing — caught by the same unverifiable path already in place — matching what this project's own resolver docstring already said this case should do.

That last checkpoint (6/34, 17.6%) held two more real bugs the remaining 6 misses traced to — neither a period-resolution gap this time. First: MD&A chunking itself. CRM's live 10-K markup shows a full sentence — *"For fiscal 2026, diluted net income per share was \$7.80 as compared to diluted net income per share of \$6.36 from a year ago."* — rendered as one `<div>` made of five sibling `<span>` elements, each wrapping a differently-styled run. `extract_paragraphs` called `soup.get_text("\n")` on the whole document, which inserts its separator between every text node it walks past — including sibling *inline* elements, not just real paragraph breaks. That split this one sentence into three separate chunks, so extraction saw "\$6.36 from a year ago" completely divorced from what it was even about. The same pattern fragmented CRM's cash-flow sentence and its remaining-performance-obligation sentence — the exact source of the "\$72.4 billion" misattributed as "revenue" earlier in this section. **Fixed** by walking leaf block-level elements instead of splitting by text node, joining each one's own text with a space. Confirmed against live data: all three sentences now come back intact. Regression tests added, including one confirming genuinely separate bullet points still stay separate (the fix must not over-merge).

Second: one bad LLM response used to crash the entire run. `Coordinator.run` called the extractor with no error handling at all — a single `AdapterParseError` (a real, reproducible "the LM returned an empty or null response") propagated straight out and lost every other chunk's already-verified claims for that ticker. Confirmed the failure wasn't about the chunk's own content — retrying the identical text in isolation succeeded 3/3 times — a transient backend issue under sustained load, not bad input. **Fixed** by catching the failure, recording it as its own trace event, and continuing — deliberately not marked as checkpoint-processed, so a resumed run retries a transient failure instead of permanently treating it as "no claims."

That checkpoint's 6 remaining misses, all CRM, traced to one more real bug — not a hypothesis, confirmed directly against CRM's own XBRL data. **CRM's own 10-K tags its annual facts with `fiscal_year=2025` , while its own 10-Q tags the SAME fiscal year's quarterly facts with `fiscal_year=2026`** — for every one of revenue, EPS, cash flow, and remaining performance obligation. The annual fact (period Feb 2025–Jan 2026, filed with the 10-K) is exactly what the MD&A's "fiscal 2026" prose describes — its values match the stated claims to the dollar (\$41.525B revenue, \$7.80 EPS, \$14.996B cash flow, \$72.4B RPO) — but the resolver's fiscal-year hint matching found a real, non-empty match at the *wrong* fact: a Q3 quarterly figure that happens to carry the "correct" number purely because CRM's own annual and quarterly filings disagree with each other on what fiscal year to call it. Not a code bug reproducing a guess — a confirmed, real inconsistency in how one filer self-tags its own SEC data across filing types. **Fixed** by preferring any annual fact over a same-numbered quarterly one, tolerating the annual fact's own number being one less than the hint, before falling back to exact matching — verified against all four real CRM concepts directly, with the common case (annual and quarterly tags agreeing, e.g. NVDA, TXN) confirmed unaffected.

THE NUMBER THAT ACTUALLY MATTERS

Re-running the full check after this fix: **153 claims, 35 resolved verdicts (33 consistent, 2 inconsistent)**. CRM's inconsistent count dropped from 6 to **2** — the 4 fiscal-year-resolution misses are gone, and the remaining 2 are exactly the already-diagnosed, untouched extraction-semantics cases (an FX-impact percentage misread as growth, and one case where the company's own stated "14 percent" RPO growth doesn't match the raw XBRL-computed 21.68% — a genuine extraction/disclosure discrepancy, not a resolver issue). **Reconciler false-positive rate given correct inputs: 0/35**. The aggregate rate landed at $2/35 = 5.7\%$ (95% half-width $\pm 7.7\text{pp}$) — a large, real drop from the original 25% baseline, not a marginal one.

That FX-impact case was a real, distinct extraction bug of its own: *"Total revenues ... was positively impacted by approximately one percent in foreign currency fluctuations"* got read as a 1% growth claim for revenue, when the sentence states only currency's partial contribution to a real, separately-stated growth number — not revenue's own total movement. **Fixed** with an explicit rule in the extraction signature: a "was impacted by X%" sentence naming a contributing factor isn't a checkable claim and should be skipped.

Verifying it surfaced something bigger: the compiled extractor artifact

(`eval/optimized_extractor.json`) freezes its signature instructions at compile time. A live docstring edit has **no effect** on the real extraction agent — what every actual caller uses — until that artifact is recompiled. Confirmed directly: the same paragraph extracted correctly zero-shot right after the docstring edit, but still wrongly through the loaded, month-old compiled file. Every earlier docstring-level fix this session, including the "compared with" rule, had never actually reached production until this recompile — they happened to not matter for their own target cases because the model was already correct there without the added instruction, which was luck, not the fix working. Recompiled, added a real labeled training example, and added a generalization eval (3 synthetic "impacted by X%" sentences with different factor words and metrics) — all pass, and a spot-check confirms legitimate growth claims are unaffected.

Re-running the full check after that fix landed one more real, distinct extraction bug: *"Total gross margin increased by 0.2 percentage points."* got read as a `growth_pct` claim of 0.20% for `GrossProfit`, a dollar concept — comparing a margin-*ratio* point change against the dollar figure's own relative growth, two completely incomparable numbers (0.2 vs. the real 5.63% dollar growth). The schema had already distinguished this — "an explicit basis-point or percentage- point change" is `bps_change`, never `growth_pct` — the field description alone just wasn't strong enough to make the model reliably apply it. **Fixed** with an explicit contrastive rule. Confirmed `bps_change` is safe to route to even without full pipeline support: the verification agent never populates a resolved claim's denominator concept — a real, separate, still-open gap — but the reconciler already returns unverifiable rather than a wrong computation when that's missing, so this fix only needed to get the claim *type* right at extraction time. Added the first real `bps_change` training example in the dataset and a generalization eval (2 synthetic cases, different metrics and phrasing, both correctly routed) — recompiled, and confirmed against the real production artifact directly.

Re-running the full check after that fix left one loose end worth chasing rather than filing away as "possibly a concept mismatch": CRM's remaining performance obligation. The company's own MD&A states *"\$72.4 billion, an increase of 14 percent year-over-year"* — pulling every real RPO fact for CRM found the true comparable prior value, **\$63.4 billion** (as of January 31, 2025): $72.4 / 63.4 = +14.2\%$, matching almost exactly. But the resolver had picked **\$59.5 billion** instead — an intervening quarterly balance from nine months earlier — implying 21.7% growth. The reason: remaining performance obligation is reported as an *instant* fact, a balance as of one

date, not a duration like revenue. The very first period-selection fix in this section (prefer a same-*length* older period over the naive "next distinct period_end" pick) computes that length from the fact's own duration — which is undefined for an instant fact, so the safeguard silently, permanently evaluated empty for every balance-style concept, falling straight back to the exact unsafe pick it was built to replace. Invisible until now because no earlier real claim in this project happened to hit a checkable year-over-year instant-fact comparison. **Fixed** by defaulting an instant fact's comparison to the closest distinct prior instant approximately one year before current's own date — the same arithmetic already used for the hinted "a year ago" path, just applied unconditionally as the default here, since an instant fact has no duration to match in the first place and MD&A almost always narrates these year-over-year anyway.

Re-running the full check after this fix left one loose end: CSCO's diluted-EPS case (19% claimed vs. 0.39% real). The chunking fix had already merged the period-establishing sentence and the EPS sentence into one chunk — extraction *could* correctly propagate the established period to the later claim, but didn't reliably: sometimes it attached the right period (correctly routed to unverifiable, since CSCO has no discrete Q4 fact), sometimes it dropped it and produced a false positive. Genuine model variance, not a deterministic bug.

The obvious fix — an explicit rule that an established period carries forward to every later claim — worked for CSCO on the first try, and broke something else on the same try: a previously-correct TXN sentence stopped expanding its dollar amounts (17.68 instead of 17,680,000,000), reproduced 3/3 with the new instruction and 3/3 correct without it — direct causal proof, not a guess. Adding a training example for the broken TXN sentence didn't help either, which led to a real, separate discovery: the compiler that bakes few-shot demonstrations into the extractor stops at the *first four* successful examples it finds walking the training set from the start — meaning **none of this session's new training examples had ever actually been used as demonstrations**. Every extraction fix this entire session worked through the signature's own instructions text alone, never through new demos. The real fix was shortening the period-propagation rule to one plain sentence, which resolved the interference — confirmed 3/3 on both the CSCO and TXN cases together.

Verifying also turned up a second, unrelated cause behind the same symptom: this CSCO paragraph produces 9 real claims, and its extraction call was regularly hitting the model's 4,000-token output limit — sometimes crashing the parse outright, sometimes truncating just short enough to drop the very last claim's period field.

Fixed by raising the limit to 6,000, confirmed directly: the same 9-claim chunk that failed or dropped fields at 4,000 completed correctly 4 times out of 4 at 6,000.

Re-running the full check after both fixes left the pipeline at zero observed false positives, but one gap still sat quietly underneath it: `bps_change` claims — a margin's own point-change, like "Total gross margin increased by 0.2 percentage points" — always came back unverifiable, not because the reconciler couldn't check a ratio, but because nothing ever told it what to divide by.

`agents/verification_agent.py` built every ratio claim without a denominator concept at all. Not a bug producing wrong verdicts — the reconciler's own fallback was already safe by construction — but a real, closeable gap in coverage.

Investigated which metric texts were actually safe to wire before adding any. "Gross margin" checked out against real CSCO numbers: computed change +20.6bps against the company's own claimed +20bps, well within tolerance. "Operating margin" looked like the identical pattern and was tested the same way — and rejected. Real CSCO GAAP operating income over revenue moved *-188bps* while the company's own MD&A states the margin "increased by 1.8 percentage points," and the sentence's own cited drivers — lower amortization, lower restructuring charges — are exactly what a non-GAAP-adjusted operating margin excludes. Wiring that mapping would have manufactured a brand-new false positive, the identical mistake already made and undone once this session for gross margin itself mapped to a dollar concept. **Fixed** narrowly: gross-margin-family metrics only, kept in their own table, deliberately smaller than the concept dictionary everything else uses — anything without an entry stays unverifiable exactly as before.

THE NUMBER THAT ACTUALLY MATTERS, FINAL

Re-running the full check after this fix: **152 claims, 33 resolved verdicts, 0 inconsistent**. Every ticker is clean. CSCO's gross-margin claim, previously unverifiable, is now correctly verified consistent — one more claim genuinely checked, zero new false positives. **Reconciler false-positive rate given correct inputs: 0/33**. The aggregate rate landed at $0/33 = 0.0\%$, holding at zero across eleven confirmed root causes — from the original 25% baseline.

A fifteenth bug turned up narrowing this control set to one vertical rather than expanding it. INTU returned zero claims — not assumed to be the same coverage gap as the financials/consumer-staples batch, checked directly. Its real heading reads *"ITEM 7 - MANAGEMENT'S DISCUSSION AND ANALYSIS..."* — dash-separated, no period — and the parser's regex only tolerated an optional period.

Fixed, with regression tests from the real heading text: INTU now returns 280 real MD&A chunks instead of 0. IBM's zero claims turned out not to be a bug at all — its 10-K states plainly *"Refer to pages 6 through 38 of IBM's 2025 Annual Report to Stockholders, which are incorporated herein by reference"* — the MD&A simply isn't in this document, a real and legitimate filing pattern, not something a heading-regex fix could reach.

0/33 proves the fixes above work — every one of them was diagnosed against a real bug found *in this same 8-ticker control set*. It doesn't prove the approach generalizes to companies never seen before. Tested directly:

`research/specificity_check_fresh.py` reruns the identical methodology against a second, disjoint set of large tech companies untouched by any fix, test, or diagnostic call this project has made — NOW, PANW, CRWD, WDAY, SNPS, FTNT, INTC, MU — also disjoint from AAPL/MSFT/NVDA/GOOGL/AMZN, the pipeline's own original dev/tuning companies.

67 claims, fresh 8-ticker set	11 resolved verdicts	5 inconsistent, all SNPS	45.5% apparent false-positive rate
---	--	--	--

A real, non-zero result — investigated rather than waved off as noise. All 5 traced to SNPS (Synopsys), splitting into two distinct causes. **Cause A:** a segment/acquisition-qualified sub-figure mislabeled under the bare metric name. SNPS's real 10-K states *"Revenues were \$7.1 billion, an increase of \$926.8 million or 15%, which includes revenues from Ansys of \$756.6 million"* — extraction grabbed the Ansys-specific contribution (\$756.6M, from SNPS's recent Ansys merger) and labeled it plain "revenue," comparing it against the company's real total revenue for an apparent 89% miss that isn't a mismatch at all, just a mislabeled number. The same sentence's China-specific decline hit the identical pattern with the sign flipped too. **Cause B:** a GAAP/non-GAAP ambiguity hiding inside an already-trusted dictionary entry. SNPS states, word for word, *"Operating income was \$1.4 billion, an increase of \$82.5 million or 6%"* — extraction read this correctly, but real GAAP `OperatingIncomeLoss` actually **fell** 32.5%, almost certainly because SNPS's stated figure excludes one-time Ansys-merger acquisition costs. The same trap already found and avoided for CSCO's "operating margin" earlier in this section — except here it hides inside `"operating income"` → `OperatingIncomeLoss`, a mapping already used successfully by other companies in the original control set. SNPS's

unusually large one-time charges just made the gap large enough to surface on a plain, previously-safe metric text.

Fixed Cause A only, by deliberate scope decision — it's bounded and matches the exact pattern of prior fixes in this section: a docstring rule teaching the extraction signature that a segment/product/subsidiary/geography-qualified sub-figure gets its *own* metric text, never the bare metric name — the resolver's existing exact-match dictionary then automatically treats the qualified text as unverifiable, no new resolver code needed. Cause B was left deliberately open: a genuinely structural question (a company's plain "operating income" prose can mean GAAP or non-GAAP depending on the company and year, and an exact-match dictionary has no way to know which), documented rather than guessed at — the same discipline as the IBM incorporated-by-reference case earlier in this document.

73 claims, re-run	9 resolved verdicts	3 inconsistent — same case, unchanged	33.3% apparent false-positive rate
---	---	---	--

HONEST SCOPE

n=9 resolved verdicts is a genuinely tiny sample — the 95% confidence half-width on 33.3% is roughly ± 31 percentage points, wide enough that the *rate* itself shouldn't be over-read. What's precise, not noisy, is the categorical finding: Cause A is confirmed fully eliminated by direct re-measurement — zero revenue-misattribution false positives remain — and every remaining miss traces to the one already-diagnosed, deliberately unfixed GAAP/non-GAAP case, not a mystery. Full writeup: docs/fresh-holdout-specificity-results.md .

08

What would full coverage actually cost?

The README's "why now" argument is directional — falling inference cost makes full coverage a cost problem, not a capability one. Turned into a real number, not a claim. Verification is free (pure Python arithmetic against already-fetched XBRL facts); retrieval is free (plain HTTP GETs); extraction is the only step with a real dollar cost, so it's the only one measured.

1,537 avg prompt tokens / call, measured	454 avg completion tokens / call, measured	215 chunks, MSFT's real 10-K MD&A	\$11.31 all 500 S&P 500 companies, one year
--	--	---	---

Real, cache-disabled extraction calls (a unique nonce per chunk rules out a stale cached response standing in for a genuine one) against MSFT's real MD&A, at `gpt-oss-20b`'s published API pricing — the exact model this project's own endpoint runs — \$0.030/1M input, \$0.130/1M output tokens, checked directly rather than assumed from memory. Eleven dollars a year for extraction across every S&P 500 company's full annual disclosure — the whole index, not one company. The prompt cost is dominated by fixed instruction overhead (~1,500 tokens), not chunk size, so cost scales with chunk *count*, not filing length, in this range.

WHAT THIS NUMBER DOESN'T SAY

It's inference cost at the current, un-fine-tuned model — not a claim about accuracy or coverage. Phase 7's own results (§02) show real reasoning gaps a zero-shot model has; §07's specificity check and `docs/robustness-and-scope.md` show real coverage gaps in what fraction of real claims resolve today. \$11.31 buys running the pipeline over every chunk, not verifying every claim a filing makes — the honest claim this number supports is narrower than "full coverage is solved."

09 Reflection

The capstone brief names this explicitly, as part of the final presentation: "a candid reflection on what you learned and what you would do differently." Most of that honesty is already scattered through the sections above, as it happened — collected here in one place rather than left implicit.

Reasoning beat optimization, and it wasn't close. Switching the extraction signature to `ChainOfThought` closed nearly the entire gap to a strong extractor; `BootstrapFewShot` and GEPA on top of that added deltas smaller than their own noise floors, and on the fresh AMZN/AAPL holdout, optimization actually made things worse. Had the whole time budget gone to reasoning-prompt design and none to optimizer runs, the result would likely look the same.

Testing for overfitting is not optional busywork — it finds real bugs. The clean 0/33 specificity result completely hid a real, fixable segment-misattribution bug that only the fresh-holdout test in §07 surfaced. A result measured only against the companies it was tuned on isn't evidence of anything beyond those companies.

A near-zero-false-positive checker is worth more than a stronger extractor. The maker/checker math in §06 isn't just theory here — nearly every real bug found across this whole project was upstream of the Reconciler, never in its own arithmetic. Investing in the checker's precision paid off more than investing further in extraction would have.

Naming a constraint openly beats working around it silently. The GRPO group size (6, not the capstone brief's suggested 8 — see `phase7/train_grpo.py`) and the GPU (a 24GB RTX 4090, not the 32GB card the brief and the original planning doc assumed) are both real deviations from plan, documented in code comments rather than smoothed over. The DSPy compiled-artifact freeze (finding 07-adjacent, discussed in §07) was the opposite lesson learned the hard way: an undocumented, invisible constraint that cost real time before anyone even noticed it was there.

WHAT WOULD BE DONE DIFFERENTLY, WITH MORE TIME

Segment/dimensional XBRL coverage is the largest remaining gap — most claims resolve "unverifiable" because the company-facts API this project reads has no Axis/Member dimensional breakdown at all. A proper XBRL dimensional parser (the full instance document, not just company-facts) would likely unlock most of the currently unresolvable segment claims (Azure, Data Center, XBOX revenue) rather than correctly but unsatisfyingly declining them. **GAAP/non-GAAP disambiguation** needs more than a dictionary — the CSCO and SNPS cases are a real starting training set for a small classifier or a second-pass check purpose-built for this one ambiguity, rather than another dictionary entry that would just be wrong for the next company. **A frontier-model comparison point** is still missing — every result in this project, including the naive-baseline ablation in §06, was measured against exactly one self-hosted `gpt-oss-20b` endpoint, so "is structure necessary" and "is structure necessary for a model this size" remain conflated. And the GRPO group size was never actually A/B tested against the brief's suggested G=8 — access to a larger card would isolate whether group size was ever the real reason `bps_change` stayed stuck, rather than leaving that as an open, unconfirmed hypothesis.

Written up as part of an ongoing capstone build. Every requirement, phase, and hardening pass in one reference: **the field guide**.

View the repository →